

# COL341

## Assignment 1 - Report

### Eshan Jain

### 2020CS50424

#### 3.1 Basic Implementation

i) Learning rate = 0.1

For the learning rate of 0.1, the gradient descent algorithm does not converge with increasing iterations. For the given data, this is too large a learning rate to effectively take small enough steps in the direction of the steepest descent in the cost function. Therefore for both the training and validation data, the MSE and MAE scores increase with the number of iterations of the gradient descent algorithm. Changing the relative tolerance parameter does not really make sense for this, as the difference in increasing with the number of iterations.

After just 10 iterations of the gradient descent algorithm, the values of the MSE and MAE scores for the training and validation data are:

Linear Regression Training MAE: 1.4289700304930604e+16  
Linear Regression Training MSE: 2.1069183181566595e+32  
Linear Regression Validation MAE: 7.520094794824873e+17

Linear Regression Validation MSE: 5.764661649009649e+35

And for 200 iterations, the MAE and MSE scores for both training and validation data overflow the float limit and are undefined.

## ii) Learning Rate = 0.01

Again for the learning rate of 0.01, the gradient descent algorithm does not converge. For the given data, this is still too large a learning rate to effectively take small enough steps in the direction of the steepest descent in the cost function. Therefore for both the training and validation data, the MSE and MAE scores increase with the number of iterations of the gradient descent algorithm. But this time, the values are not as large as when the learning rate was 0.1.

After 10 iterations, we have the following:

Linear Regression Training MAE: 2557949.5798920966

Linear Regression Training MSE: 6751269111006.135

Linear Regression Validation MAE: 11119304.065800318

Linear Regression Validation MSE: 126032457599460.77

After 200 iterations, overflow has not occurred yet, and the values are:

Linear Regression Training MAE: 1.744605461532271e+126

Linear Regression Training MSE: 3.1404791427653817e+252

Linear Regression Validation MAE: 7.583729801570534e+126

Linear Regression Validation MSE: 5.86263550915762e+253

However, 500 iterations do cause an overflow in both the MAE and MSE scores.

iii) Learning Rate = 0.001

This is a good enough learning rate for the given data, and the gradient descent algorithm converges in this case.

In comparison to the above two learning rates, we have:

After 10 iterations, we have the following:

Linear Regression Training MAE: 0.6598852183034914  
Linear Regression Training MSE: 0.7043224286023175  
Linear Regression Validation MAE: 0.6889660970549375  
Linear Regression Validation MSE: 0.7314702265878003

After 200 iterations, we have the following:

Linear Regression Training MAE: 0.5101726130974069  
Linear Regression Training MSE: 0.40633483214109956  
Linear Regression Validation MAE: 0.5664843867152056  
Linear Regression Validation MSE: 0.5641457987102589

After setting the max iterations to 1000 and the relative tolerance to 0.001, we have the following scores:

Linear Regression Training MAE: 0.4831120120021878  
Linear Regression Training MSE: 0.35640642214944795  
Linear Regression Validation MAE: 0.5589828050663396  
Linear Regression Validation MSE: 0.5663776806392639

Now keeping the max iterations constant and changing the relative tolerance to 0.0001, we have:

Linear Regression Training MAE: 0.3761420631427545  
Linear Regression Training MSE: 0.22294502987972967  
Linear Regression Validation MAE: 0.6234446184856276  
Linear Regression Validation MSE: 0.6449158474671899

We observe that the model is now going towards overfitting as on the training set, we achieve a better performance, but it does not generalize well, and we see a worse performance on the validation set.

Now keeping the max iterations constant and changing the relative tolerance to 0.01, we have:

Linear Regression Training MAE: 0.6333912305340261  
Linear Regression Training MSE: 0.6577928864063601  
Linear Regression Validation MAE: 0.6559258831892438  
Linear Regression Validation MSE: 0.6773315133825653

Here we observe underfitting, as the model gives worse performance for both the training and validation set since the threshold is quite weak.

### 3.2 Ridge Regression

We have the cost function for ridge regression as:-

$$E_{in}(w) = \frac{1}{M} (\sum_{i=1}^M (y_i - \sum_{j=1}^N \theta_j x_{ij})^2 + \lambda \sum_{j=1}^N w_j^2)$$

Now since the first term is same as that in linear regression, only the second term would cause a change in the parameter update equation.

The gradient of the second term would be :-

$$\frac{1}{M} \lambda \theta$$

Therefore the weight update equation for ridge regression would be:-

$$\bar{w} = w - \alpha \left( \frac{1}{M} [X^T (X\theta - y) + \lambda \theta] \right)$$

Where the first term comes from the linear regression part and the second term from the regularisation part.

i)  $\lambda = 5$

Again for the learning rates of 0.1 and 0.01, the gradient descent is diverging and, therefore, the MAE and MSE scores are undefined.

However, in the case when the learning rate is 0.001, we have the following:

After 10 iterations;

Ridge Regression Training MAE: 0.6598957872458188

Ridge Regression Training MSE: 0.7043420244013787

Ridge Regression Validation MAE: 0.6889752577115832

Ridge Regression Validation MSE: 0.7314937923262808

After 200 iterations:

Ridge Regression Training MAE: 0.5103790558737039

Ridge Regression Training MSE: 0.40680298019138417

Ridge Regression Validation MAE: 0.5665709367679435

Ridge Regression Validation MSE: 0.5641384220562985

After 1000 iterations and the relative tolerance of 0.001 :

Ridge Regression Training MAE: 0.48434238917139893

Ridge Regression Training MSE: 0.3585660265747396

Ridge Regression Validation MAE: 0.5588074665117294

Ridge Regression Validation MSE: 0.5659436777520362

Therefore we observe slightly better performance in ridge regression as compared to the linear regression, although a significant change is not observed due to the small value of the regularisation parameter  $\lambda$ .

ii)  $\lambda = 25$

Now again, for the values of learning rate 0.1 and 0.01, we don't see the gradient descent converging.

In the case when the learning rate is 0.001, we observe the following:

After 10 iterations:-

Ridge Regression Training MAE: 0.6599380463195423

Ridge Regression Training MSE: 0.7044229194245482

Ridge Regression Validation MAE: 0.6890118808769022

Ridge Regression Validation MSE: 0.7315905471335976

After 200 iterations

Ridge Regression Training MAE: 0.5111945681132539

Ridge Regression Training MSE: 0.40867517139613574

Ridge Regression Validation MAE: 0.566913262021991

Ridge Regression Validation MSE: 0.5641320240341328

After 1000 iterations:

Ridge Regression Training MAE: 0.48871660015167134

Ridge Regression Training MSE: 0.366471793778875

Ridge Regression Validation MAE: 0.5586344377318001

Ridge Regression Validation MSE: 0.5645936705730505

Therefore we observe a much better performance with  $\lambda = 25$ , as compared to  $\lambda = 5$ . This is because as we penalize the model complexity more, it prevents overfitting, and although we achieve a slightly worse performance on the training set, we are able to generalize better and so achieve a greater performance on the validation set.



### 3.3 Using Scikit-Learn Library

The performance of sklearn models on the validation dataset are:

Scikitlearn Linear Regression MAE: 0.8322387027130557

Scikitlearn Linear Regression MSE: 1.025165437413566

Scikitlearn Ridge Regression MAE: 0.7803008143298173

Scikitlearn Ridge Regression MSE: 0.9319081640199326

Which is quite poor compared to our model's performance.

### 3.4 Feature Selection

#### i) SelectKBest

Trained the linear regression model by selecting the best 10 features using the SelectKBest, and here are the results:-

For learning rate = 0.1:

After 10 iterations:

Linear Regression Training MAE: 2.3669448302135216

Linear Regression Training MSE: 7.020847634794508

After 1000 iterations:

Linear Regression Training MAE: 0.8875548211429107

Linear Regression Training MSE: 1.1421340055410378

A significant difference is that now the model converges even for learning rate = 0.1. This is because now, for the reduced number of features, this value of the learning rate is not too large, and in each iteration, the gradient descent algorithm can take small enough steps to ensure that we are moving in the right direction and toward convergence. So these MSE and MAE scores are much better than the undefined values in the case of the linear regression model trained on all features.

For the learning rate = 0.01:

After 10 iterations:

Linear Regression Training MAE: 4.819840559488586

Linear Regression Training MSE: 25.765714158705947

After 1000 iterations:

Linear Regression Training MAE: 0.9173913833914438

Linear Regression Training MSE: 1.2350872636970207

For the learning rate = 0.001:

After 10 iterations:

Linear Regression Training MAE: 5.290724008794868

Linear Regression Training MSE: 30.25872745594136

After 1000 iterations:

Linear Regression Training MAE: 2.266571104658656

Linear Regression Training MSE: 6.500151949323961

Therefore we observe that with the reduced number of features, the learning rate of 0.1 produces the best results, followed by 0.01 and 0.001, which is in stark contrast to the model trained on all features, which followed the opposite order.

## ii) SelectFromModel

Trained the ridge regression model by selecting the best 10 features using the SelectFromModel, and here are the results:-

For lambda = 5:

For learning rate = 0.1,

After 10 iterations;

Ridge Regression Training MAE: 0.774502673098934

Ridge Regression Training MSE: 1.0795864757287257

After 1000 iterations:

Ridge Regression Training MAE: 0.6204450830121448

Ridge Regression Training MSE: 0.6263751375381954

For learning rate = 0.001,

After 10 iterations:

Ridge Regression Training MAE: 5.065274028807056

Ridge Regression Training MSE: 27.639600534427284

After 1000 iterations:

Ridge Regression Training MAE: 0.8666754758268737

Ridge Regression Training MSE: 1.2963249086288018

For  $\lambda = 25$ :

For learning rate = 0.1,

After 10 iterations;

Ridge Regression Training MAE: 0.7931300211820306

Ridge Regression Training MSE: 1.115196739057996

After 1000 iterations:

Ridge Regression Training MAE: 0.6722545032500185

Ridge Regression Training MSE: 0.7692863063211384

For learning rate = 0.001,

After 10 iterations:

Ridge Regression Training MAE: 5.0654299434679775

Ridge Regression Training MSE: 27.641308799926264

After 1000 iterations:

Ridge Regression Training MAE: 0.8960875825019998

Ridge Regression Training MSE: 1.3458784844253513

Similar to what we observed in the linear regression model, with a reduced number of features, we get better results for the learning rate of 0.1 than 0.001. Also, as expected, we see better results with  $\lambda = 25$  than with  $\lambda = 5$ , as was the case in the ridge regression model trained on all the features.

### 3.5 Classification

The cost function for each class is:

$$J(\theta) = \frac{1}{M} \sum_{i=1}^M (y_i \log(h_{\theta}(x_i)) + (1 - y_i) \log(1 - h_{\theta}(x_i)))$$


For each class, we create a binary output from the given output, which reflects whether the output is that class or not for all data elements. We update the weight vector based on this modified output and repeat this for each class separately. Therefore the weight update equation is:-

$$\theta_j = \theta_j - \alpha \left( \frac{1}{M} \sum_{i=1}^M ((h_{\theta}(x_i) - y_i) x_{ji}) \right)$$

## 3.6 Visualization

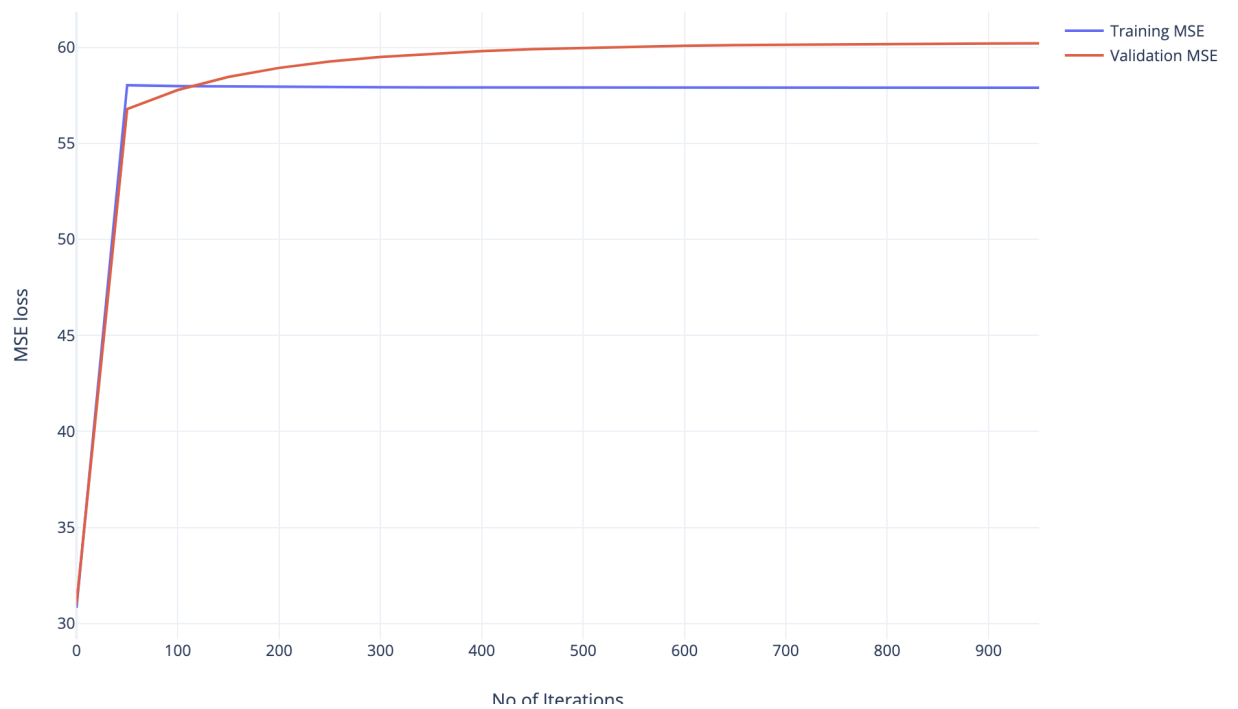
3.1 Without normalising, only the learning rate of 0.001 converges.

For learning rate = 0.001:




### 3.1 Using Normalised Gradient

For learning rate = 0.1:






For learning rate = 0.01:



For learning rate = 0.001:




## 3.2

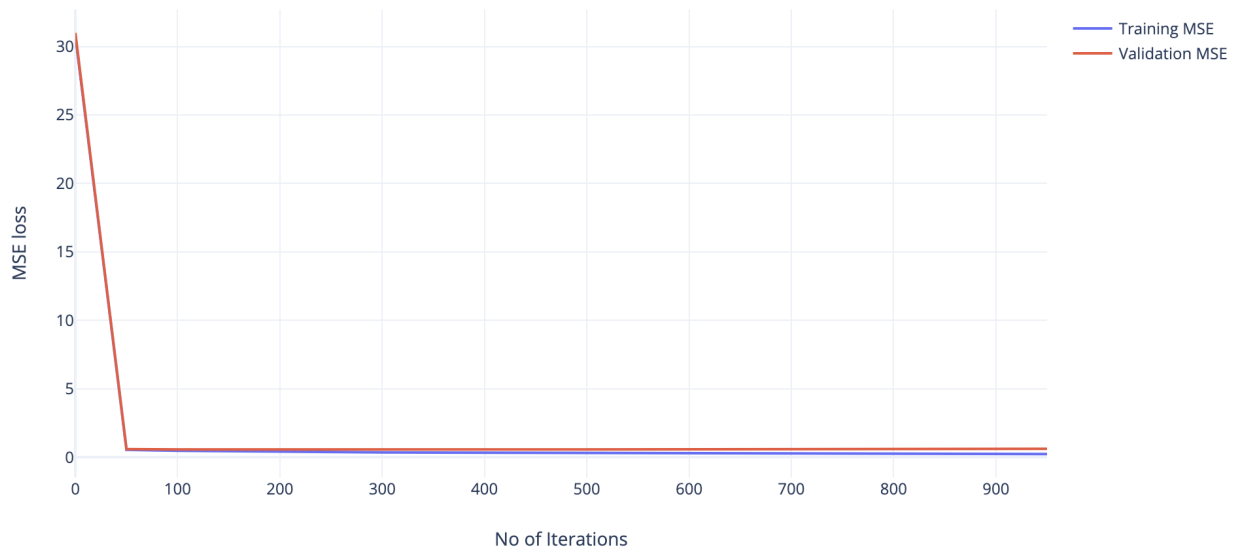
As mentioned, the gradient descent converges only when the learning rate = 0.001.

So for the learning rate = 0.001:

$\text{Lambda} = 5$



Lambda = 25




As expected, we observe better results with  $\lambda = 25$ , since a higher penalty for the weights prevents overfitting, as can be seen from the plots (better performance on the validation set for  $\lambda = 25$ ).

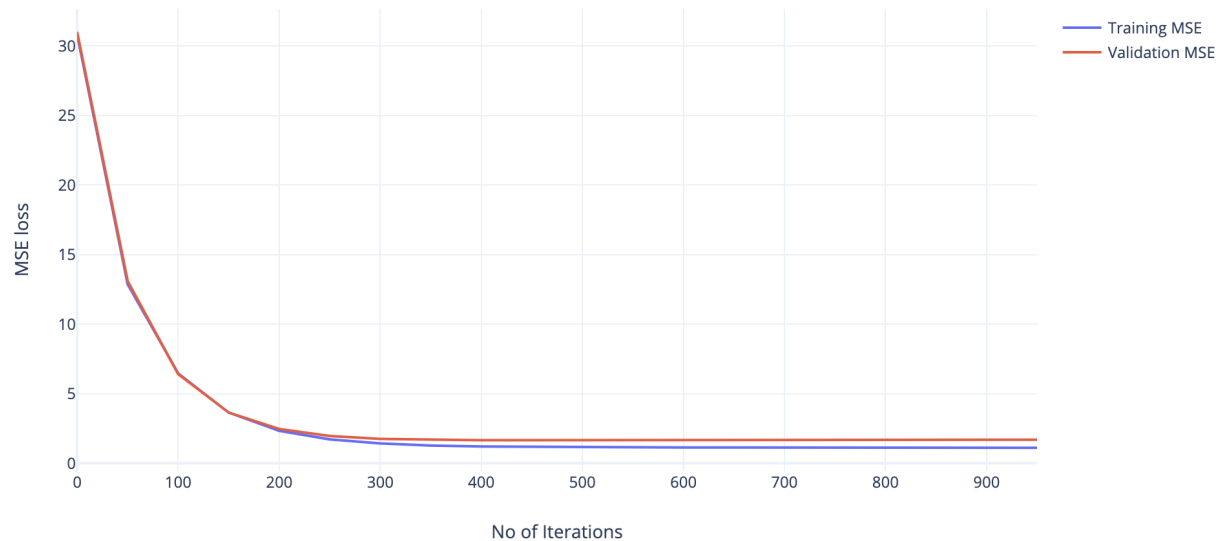
### 3.4

Linear Regression with top 10 features:

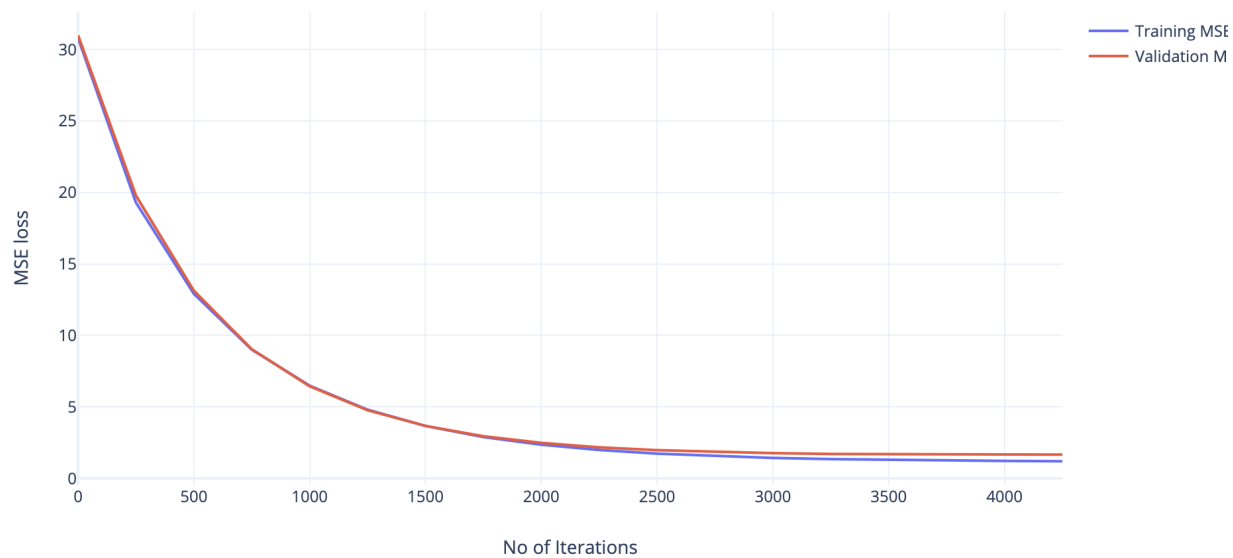
Learning Rate = 0.1



Learning Rate = 0.01




Learning Rate = 0.001




Ridge Regression using top 10 features:

Learning Rate = 0.1

Lambda = 5




Lambda = 25




Learning Rate = 0.001

Lamda = 5





Lambda = 25



Again, the plots show that ridge regression prevents overfitting as the validation error is reduced as compared to the plots of linear regression.

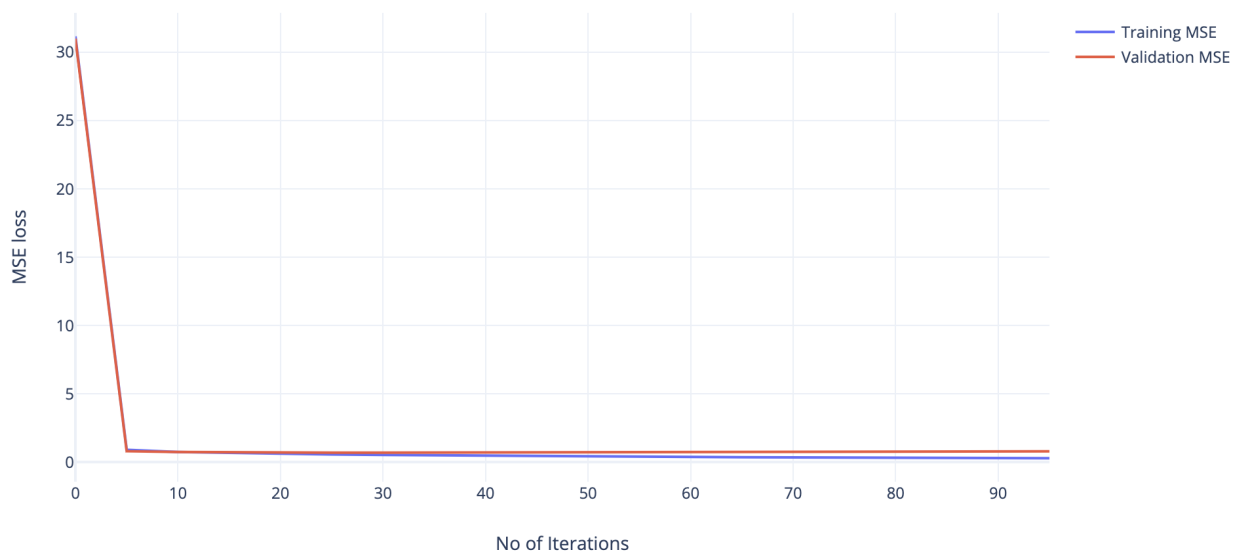
3.5

## Linear regression trained on parts of training dataset

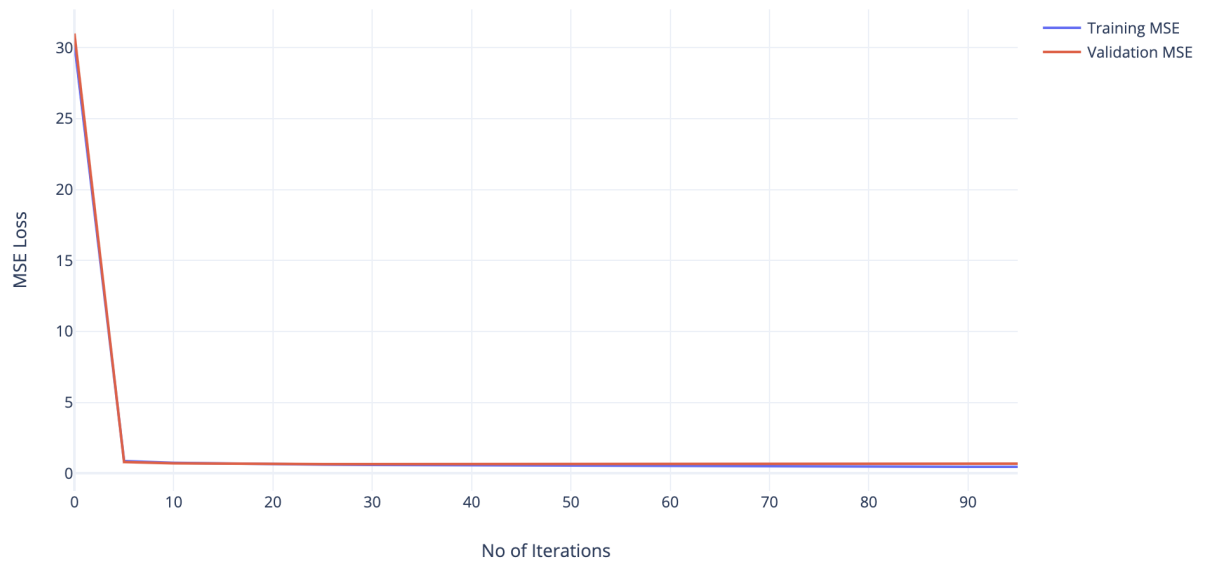
For 3.1:

Learning rate = 0.001 (since it only converges for this)


One-fourth of training data:-




One half of training data:



Three Fourth of training data:



Full Data:



### **3.1 Linear Regression on two halves of training data:**

Mean Absolute Difference = 1.3851158525138192

### **3.2 Ridge Regression on two halves of training data:**

For lambda = 5:

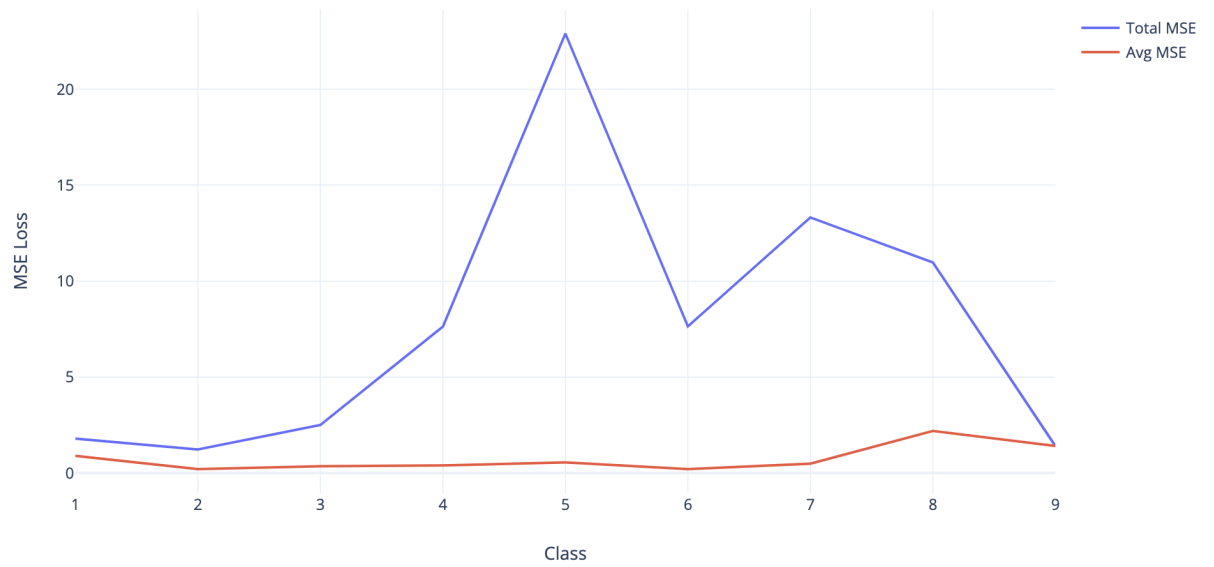
Mean Absolute Difference= 1.384426093447689

For lambda = 25

Mean Absolute Difference= 1.3816830903902242


Ridge regression clearly shows better performance as expected.

### **3.1 Total and Avg MSE Loss per class for linear regression**




### 3.4 SelectFromModel Feature selection for 10,100, 1000, 2048 features

For learning rate = 0.001, we have:



**3.7**



This is the plot for  $E_{in} - E_{out}$  vs the number of dimensions (2, 5, 10, 100).



